

# PD1 full: LOO over 23 tasks, neglog warp (115 runs): regret vs wall-clock

per task (leave-one-out: every task trains its own prior)

